

Advancing EEG-Based Diagnostics: A Comprehensive Data Analysis Approach for the Classification of Neurological Patterns using Machine Learning



IE6400 – FOUNDATIONS OF DATA ANALYTICS

Project Report 3

Group Number 17

Janani Karthikeyan (002830003)

Milan Gurumurthy (002833029)

Prathyusha Adhikam (002835277)

Saathvika Kethineni (002893814)

Shreyas Sreenivas (002825934)

Contents:

1. Abstract.....
2. Introduction and Background Information.....
3. Data Sources.....
4. Methods and Results.....
5. Summary and Conclusions

Abstract:

Epilepsy, which affects millions of people worldwide, necessitates precise and timely diagnosis in order to provide effective treatment and risk reduction. Manual analysis of electroencephalogram (EEG) recordings presents difficulties due to time constraints and expert variability. This study is the first to use Convolutional Neural Networks (CNNs) to directly process raw EEG signals, eliminating the need for manual feature extraction. We compare the performance of time and frequency domain signals in detecting epileptic activity using the intracranial Freiburg and scalp CHB-MIT databases. Our three-tiered classification experiments, which include binary and ternary scenarios, show that frequency domain signals have superior accuracy, emphasizing their potential for CNN applications.

Introduction and Background Information:

Epilepsy, the world's second most common neurological disorder, necessitates effective diagnostic methods in order to initiate timely interventions. Manual EEG analysis is time-consuming and prone to interpretational discrepancies among experts in current practice. To address this, our study is the first to use a machine learning approach, specifically Convolutional Neural Networks (CNNs), to directly process raw EEG signals for improved seizure detection.

Traditional diagnostic methods involve complex feature extraction and classification techniques that frequently necessitate manual intervention. Deep learning, particularly CNNs, has emerged as a promising means of speeding up this process. Our research focuses on the timely and accurate identification of epileptic patterns in EEG recordings, which is critical for initiating antiepileptic drug therapy and lowering associated risks.

Previous research in epilepsy diagnosis focused primarily on feature extraction,

which was frequently done in the time domain. Recent advances in deep learning, on the other hand, have demonstrated the utility of using raw signals for classification tasks. In this context, our research delves into the frequency domain, investigating hidden information within EEG signals and comparing its effectiveness to time domain signals.

The goal of this project is to create a robust classification model capable of analyzing EEG data and classifying it into distinct classes. This endeavor is especially important in the fields of neuroscience and medicine, where EEG data is used to diagnose neurological disorders such as epilepsy.

To address these issues, this project will conduct a thorough examination of EEG data, utilizing two distinct datasets: the CHB-MIT EEG Database, which includes various seizure types as well as non-seizure data, and the Bonn EEG Dataset, which focuses specifically on epileptic seizures. The project roadmap tasks cover critical steps in the data analysis pipeline, such as preprocessing and feature extraction, as well as model selection, training, and evaluation.

Our study found that using frequency domain signals in our experiments led to much higher accuracy in classifying epileptic patterns compared to using time domain signals. The Convolutional Neural Network (CNN), known for its effectiveness in recognizing images and videos, proves to be a powerful tool for analyzing EEG data, which is complex and high-dimensional. In essence, our research supports a shift in how we approach epilepsy diagnosis—favoring a data-driven and automated method that not only improves accuracy but also speeds up the diagnostic process.

Data Sources:

<https://physionet.org/content/chbmit/1.0.0/>

<https://www.ukbonn.de/epileptologie/arbeitsgruppen/ag-lehnertz-neurophysik/dow>

[nloads/](#)

Our project draws on two critical EEG datasets, each of which contributes unique perspectives to our analysis. The CHB-MIT EEG Database contains a large collection of EEG recordings from people who have epilepsy. This dataset includes both seizure and non-seizure data, resulting in a diverse and clinically relevant range of neurological patterns.

In addition, the Bonn EEG Dataset focuses on epileptic seizures, providing a concentrated set of recordings for in-depth examination. These datasets form the foundation of our study, allowing us to train and evaluate our classification model on a diverse set of EEG data, which is critical for furthering our understanding of neurological conditions, particularly epilepsy.

Methods and Results:

Tasks:

Data Preprocessing and Feature Extraction:

Data preprocessing requires downloading and extracting datasets, exploring their structure, and carrying out necessary steps such as missing value handling, noise reduction, and data augmentation. The extraction of relevant features from EEG signals, taking into account both time-domain and frequency-domain features, becomes a critical step. Data splitting then ensures the creation of training, validation, and test sets, laying the groundwork for model development.

```
In [42]: 1 # Import necessary Libraries
2 from glob import glob
3 import os
4 import mne
5 import numpy as np
6 import matplotlib.pyplot as plt
7 import pandas as pd
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10 from sklearn.metrics import confusion_matrix, precision_score, recall_score
11 from scipy import stats
12 import numpy as np
13 from tqdm import tqdm_notebook
14 from sklearn.linear_model import LogisticRegression
15 from sklearn.pipeline import Pipeline
16 from sklearn.preprocessing import StandardScaler
17 from sklearn.model_selection import GroupKFold, GridSearchCV, cross_val_score
18 from tensorflow.keras.layers import Conv1D, BatchNormalization, LeakyReLU, MaxPooling1D, Dense
19 from tensorflow.keras.models import Sequential
20 from tensorflow.keras.backend import clear_session
21 from sklearn.model_selection import GroupKFold, LeaveOneGroupOut
22 from sklearn.preprocessing import StandardScaler
```

```
In [43]: 1 # Read all files ending with .edf and .edf.seizures
2 all_files_path = glob(r"D:\Documents\Prof_Docs\FDA\Projects\Project 3\chb01\*.edf")
3           glob(r"D:\Documents\Prof_Docs\FDA\Projects\Project 3\chb01\*.edf.seizures")
4
5
6 print("Total files:", len(all_files_path))
```

Total files: 23

```
In [44]: 1 all_files_path[0]
```

```
Out[44]: 'D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb01\\chb01_03.edf'
```

In [45]:

```
1 # Segregate files with seizures and their corresponding EEG data files
2 files_with_seizures = []
3 eeg_files_with_seizures = []
4
5 # Create a copy of the all_files_path list
6 all_file_path = all_files_path.copy()
7
8 for seizure_file_path in all_files_path:
9     if '.edf.seizures' in seizure_file_path:
10         # This is a seizure data file, so add it to the list
11         files_with_seizures.append(seizure_file_path)
12
13     # Extract the corresponding EEG data file name
14     eeg_file_name = os.path.basename(seizure_file_path).replace('.edf')
15     eeg_file_path = next((path for path in all_files_path if eeg_file_
16
17         if eeg_file_path:
18             # Remove the corresponding EEG data file from the copy
19             all_file_path.remove(eeg_file_path)
20             eeg_files_with_seizures.append(eeg_file_path)
21
22     # Remove files with '.edf.seizures' extension from the remaining files
23     healthy_file_path = [path for path in all_file_path if '.edf.seizures' not in
24
25     # Print or use the segregated file paths
26     print("EEG files with seizures corresponding to seizure files:", len(eeg_
27     print(eeg_files_with_seizures)
```

```
EEG files with seizures corresponding to seizure files: 7
['D:\\\\Documents\\\\Prof_Docs\\\\FDA\\\\Projects\\\\Project 3\\\\chb-mit-scalp-eeg-database-1.0.0\\\\chb01\\\\chb01_03.edf', 'D:\\\\Documents\\\\Prof_Docs\\\\FDA\\\\Projects\\\\Project 3\\\\chb-mit-scalp-eeg-database-1.0.0\\\\chb01\\\\chb01_04.edf', 'D:\\\\Documents\\\\Prof_Docs\\\\FDA\\\\Projects\\\\Project 3\\\\chb-mit-scalp-eeg-database-1.0.0\\\\chb01\\\\chb01_15.edf', 'D:\\\\Documents\\\\Prof_Docs\\\\FDA\\\\Projects\\\\Project 3\\\\chb-mit-scalp-eeg-database-1.0.0\\\\chb01\\\\chb01_16.edf', 'D:\\\\Documents\\\\Prof_Docs\\\\FDA\\\\Projects\\\\Project 3\\\\chb-mit-scalp-eeg-database-1.0.0\\\\chb01\\\\chb01_18.edf', 'D:\\\\Documents\\\\Prof_Docs\\\\FDA\\\\Projects\\\\Project 3\\\\chb-mit-scalp-eeg-database-1.0.0\\\\chb01\\\\chb01_21.edf', 'D:\\\\Documents\\\\Prof_Docs\\\\FDA\\\\Projects\\\\Project 3\\\\chb-mit-scalp-eeg-database-1.0.0\\\\chb01\\\\chb01_26.edf']
```

, 5:06 PM

EEG classification_Final - Jupyter Notebook

In [46]:


```
1 print("Remaining files without seizures:", len(healthy_file_path))
2 print(healthy_file_path)
```

```
Remaining files without seizures: 9
['D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb-mit-scalp-eeg-database-1.0.0\\chb01\\chb01_17.edf', 'D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb-mit-scalp-eeg-database-1.0.0\\chb01\\chb01_19.edf', 'D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb-mit-scalp-eeg-database-1.0.0\\chb01\\chb01_20.edf', 'D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb-mit-scalp-eeg-database-1.0.0\\chb01\\chb01_22.edf', 'D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb-mit-scalp-eeg-database-1.0.0\\chb01\\chb01_23.edf', 'D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb-mit-scalp-eeg-database-1.0.0\\chb01\\chb01_24.edf', 'D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb-mit-scalp-eeg-database-1.0.0\\chb01\\chb01_25.edf', 'D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb-mit-scalp-eeg-database-1.0.0\\chb01\\chb01_27.edf', 'D:\\Documents\\Prof_Docs\\FDA\\Projects\\Project 3\\chb-mit-scalp-eeg-database-1.0.0\\chb01\\chb01_29.edf']
```

```

1 # Count of the files
2 seizure_files_count = len(files_with_seizures)
3 healthy_files_count = len(healthy_file_path)
4
5 # Visualization
6 plt.figure(figsize=(10, 6))
7
8 # Plotting the count of corresponding EEG files using a bar plot
9 plt.bar([0], [seizure_files_count], color='blue', label='EEG Files with Seizures')
10
11 # Plotting the count of remaining healthy files using a bar plot
12 plt.bar([1], [healthy_files_count], color='green', label='Remaining Healthy Files')
13
14 plt.xticks([0, 1], ['Seizure Files', 'Healthy Files'])
15 plt.ylabel('File Count')
16 plt.title('Count of Seizure Files vs Remaining Healthy Files')
17 plt.legend()
18 plt.show()


```



```

1 # Taking the first file from each path
2 seizure_file = eeg_files_with_seizures[0]
3 healthy_file = healthy_file_path[0]
4
5 # Reading EEG data from seizure file
6 seizure_raw = mne.io.read_raw_edf(seizure_file, preload=True)
7
8 # Reading EEG data from healthy file
9 healthy_raw = mne.io.read_raw_edf(healthy_file, preload=True)
10
11 # Plotting EEG data for one seizure file
12 fig, ax = plt.subplots(2, 1, figsize=(15, 8), sharex=True)
13 seizure_raw.plot_psd(ax=ax[0], fmin=0, fmax=50, show=False)
14 ax[0].set_title('Power Spectral Density - Seizure File')
15
16 # Plotting EEG data for one healthy file
17 healthy_raw.plot_psd(ax=ax[1], fmin=0, fmax=50, show=False)
18 ax[1].set_title('Power Spectral Density - Healthy File')
19
20 plt.tight_layout()
21 plt.show()

```



```
1 [49]: 1 # Define a function to read and preprocess EEG data from a given file
2 def read_data(file_path):
3     # Read raw EEG data from the specified file and preload it into memory
4     datax = mne.io.read_raw_edf(file_path, preload=True)
5
6     # Print the shape of the original EEG data
7     print("Original data shape:", datax.get_data().shape)
8
9     # Set EEG reference to average reference and apply bandpass filtering
10    datax.set_eeg_reference()
11    datax.filter(l_freq=1, h_freq=45)
12
13    # Print the shape of the processed EEG data
14    print("Processed data shape:", datax.get_data().shape)
15
16    # Create fixed-length epochs from the preprocessed data with a specific duration
17    epochs = mne.make_fixed_length_epochs(datax, duration=25, overlap=0)
18
19    # Print the shape of the epochs data before loading
20    print("Epochs data shape before loading:", epochs.get_data().shape)
21
22    # Load the data into memory for all epochs, drop bad epochs, and print the number of epochs created
23    epochs.load_data()
24    epochs.drop_bad()
25    print("Number of epochs created:", len(epochs))
26
27    # Get the data from the epochs and print the final shape
28    epochs_data = epochs.get_data()
29    print("Final epochs data shape:", epochs_data.shape)
30
31    # Return the preprocessed and segmented EEG data
32    return epochs_data
```

```
In [51]: 1 # Getting the shape to have an idea for our input shape  
2 datax.shape
```

```
Out[51]: (144, 23, 6400)
```

```
In [52]: 1 # Filtering all the files  
2 control_epochs_array=[read_data(subject) for subject in healthy_file_path]  
3 patients_epochs_array=[read_data(subject) for subject in eeg_files_with_se
```

```
Extracting EDF parameters from D:\Documents\Prof_Docs\FDA\Projects\Project  
3\chb-mit-scalp-eeg-database-1.0.0\chb01\chb01_17.edf...  
EDF file detected  
Setting channel info structure...  
Creating raw.info structure...  
Reading 0 ... 921599 = 0.000 ... 3599.996 secs...  
  
C:\Users\shrey\AppData\Local\Temp\ipykernel_24812\1607341177.py:4: Runtime  
Warning: Channel names are not unique, found duplicates for: {'T8-P8'}. Ap  
plying running numbers for duplicates.  
    datax = mne.io.read_raw_edf(file_path, preload=True)  
  
Original data shape: (23, 921600)  
EEG channel type selected for re-referencing  
Applying average reference.  
Applying a custom ('EEG',) reference.  
Filtering raw data in 1 contiguous segment  
Setting up band-pass filter from 1 - 45 Hz  
  
FIR filter parameters
```

```
In [53]: 1 # getting the length of the files for reference  
2 control_epochs_labels=[len(i)*[0] for i in control_epochs_array]  
3 patients_epochs_labels=[len(i)*[1] for i in patients_epochs_array]  
4 print(len(control_epochs_labels),len(patients_epochs_labels))
```

```
9 7
```

```
In [54]: 1 # Storing all files  
2 data_list=control_epochs_array+patients_epochs_array  
3 label_list=control_epochs_labels+patients_epochs_labels  
4 print(len(data_list),len(label_list))
```

```
16 16
```

```
In [55]: 1 # Grouping the files  
2 groups_list=[[i]*len(j) for i, j in enumerate(data_list)]
```

```
In [56]: 1 # Stacking the data  
2 data_array=np.vstack(data_list)  
3 label_array=np.hstack(label_list)  
4 group_array=np.hstack(groups_list)  
5 print(data_array.shape,label_array.shape,group_array.shape)
```

```
(2095, 23, 6400) (2095,) (2095,)
```

06 PM

EEG classification_Final - Jupyter Notebook


```
1 [57]: # Function to calculate the mean along the last axis of the input data
2 def mean(data):
3     return np.mean(data, axis=-1)
4
5 # Function to calculate the standard deviation along the last axis of the input data
6 def std(data):
7     return np.std(data, axis=-1)
8
9 # Function to calculate the peak-to-peak range along the last axis of the input data
10 def ptp(data):
11     return np.ptp(data, axis=-1)
12
13 # Function to calculate the variance along the last axis of the input data
14 def var(data):
15     return np.var(data, axis=-1)
16
17 # Function to find the minimum value along the last axis of the input data
18 def minim(data):
19     return np.min(data, axis=-1)
20
21 # Function to find the maximum value along the last axis of the input data
22 def maxim(data):
23     return np.max(data, axis=-1)
24
25 # Function to find the index of the minimum value along the last axis of the input data
26 def argminim(data):
27     return np.argmin(data, axis=-1)
28
29 # Function to find the index of the maximum value along the last axis of the input data
30 def argmaxim(data):
31     return np.argmax(data, axis=-1)
32
33 # Function to calculate the mean square along the last axis of the input data
34 def mean_square(data):
35     return np.mean(data**2, axis=-1)
36
37 # Function to calculate the root mean square along the last axis of the input data
38 def rms(data):
39     return np.sqrt(np.mean(data**2, axis=-1))
40
41 # Function to calculate the absolute differences between consecutive elements along the last axis of the input data
42 def abs_diffs_signal(data):
43     return np.sum(np.abs(np.diff(data, axis=-1)), axis=-1)
44
45 # Function to calculate the skewness along the last axis of the input data
46 def skewness(data):
47     return stats.skew(data, axis=-1)
48
49 # Function to calculate the kurtosis along the last axis of the input data
50 def kurtosis(data):
51     return stats.kurtosis(data, axis=-1)
52
53 # Function to concatenate a set of statistical features along the last axis of the input data
54 def concatenate_features(data):
55     return np.concatenate((mean(data), std(data), ptp(data), var(data), minim(data),
56                           argminim(data), argmaxim(data), mean_square(data)))
```

Model Selection and Training:

Model selection becomes critical, and options such as Convolutional Neural Networks (CNNs) or Recurrent Neural Networks (RNNs) are considered due to their ability to handle sequential data such as EEG signals. The following steps involve model training, in which the chosen model is trained on the training data while implementing overfitting mitigation strategies such as dropout or early stopping.

```
In [59]: 1 # Initialize a Logistic Regression classifier with a specified maximum number of iterations
2 clf = LogisticRegression(max_iter=1000)
3
4 # Initialize a GroupKFold cross-validator with 5 splits
5 gkf = GroupKFold(n_splits=5)
6
7 # Define the parameter grid for hyperparameter tuning
8 param_grid = {'classifier__C': [0.01, 0.05, 0.1, 0.5, 1, 2, 3, 4, 5, 8, 16, 32, 64, 128, 256, 512, 1024]}
9
10 # Create a pipeline with a standard scaler and the Logistic regression classifier
11 pipe = Pipeline([('scaler', StandardScaler()), ('classifier', clf)])
12
13 # Initialize GridSearchCV with the pipeline, parameter grid, GroupKFold, and n_jobs
14 gscv = GridSearchCV(pipe, param_grid, cv=gkf, n_jobs=16)
15
16 # Fit the GridSearchCV to the features, labels, and group information
17 gscv.fit(features, label_array, groups=group_array)
```

Out[59]:



```
In [60]: 1 # Getting the best score
2 gscv.best_score_
```

Out[60]: 0.7681923293852229

```

1 # Function to define a CNN model
2 def cnnmodel():
3     # Clear any previous TensorFlow sessions and models
4     clear_session()
5
6     # Create a Sequential model
7     model = Sequential()
8
9     # Convolutional Layer 1
10    model.add(Conv1D(filters=3, kernel_size=3, strides=1, input_shape=(64, 1)))
11    model.add(BatchNormalization())
12    model.add(LeakyReLU())
13    model.add(MaxPool1D(pool_size=2, strides=2)) # 2
14
15    # Convolutional Layer 2
16    model.add(Conv1D(filters=3, kernel_size=3, strides=1)) # 3
17    model.add(LeakyReLU())
18    model.add(MaxPool1D(pool_size=2, strides=2)) # 4
19    model.add(Dropout(0.3))
20
21    # Convolutional Layer 3
22    model.add(Conv1D(filters=3, kernel_size=3, strides=1)) # 5
23    model.add(LeakyReLU())
24    model.add(AveragePooling1D(pool_size=2, strides=2)) # 6
25    model.add(Dropout(0.3))
26
27    # Convolutional Layer 4
28    model.add(Conv1D(filters=3, kernel_size=3, strides=1)) # 7
29    model.add(LeakyReLU())
30    model.add(AveragePooling1D(pool_size=2, strides=2)) # 8
31
32    # Convolutional Layer 5
33    model.add(Conv1D(filters=3, kernel_size=3, strides=1)) # 9
34    model.add(LeakyReLU())
35
36    # Global Average Pooling Layer
37    model.add(GlobalAveragePooling1D()) # 10
38
39    # Output Layer
40    model.add(Dense(1, activation='sigmoid')) # 11
41
42    # Compile the model with Adam optimizer and binary crossentropy Loss
43    model.compile('adam', loss='binary_crossentropy', metrics=['accuracy'])
44
45    return model
46
47 # Create an instance of the CNN model
48 model = cnnmodel()
49
50 # Display the model summary
51 model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv1d (Conv1D)	(None, 6398, 3)	210
batch_normalization (Batch Normalization)	(None, 6398, 3)	12
leaky_re_lu (LeakyReLU)	(None, 6398, 3)	0
max_pooling1d (MaxPooling1D)	(None, 3199, 3)	0
conv1d_1 (Conv1D)	(None, 3197, 3)	30
leaky_re_lu_1 (LeakyReLU)	(None, 3197, 3)	0
max_pooling1d_1 (MaxPooling1D)	(None, 1598, 3)	0
dropout (Dropout)	(None, 1598, 3)	0
conv1d_2 (Conv1D)	(None, 1596, 3)	30
leaky_re_lu_2 (LeakyReLU)	(None, 1596, 3)	0
average_pooling1d (AveragePooling1D)	(None, 798, 3)	0
dropout_1 (Dropout)	(None, 798, 3)	0
conv1d_3 (Conv1D)	(None, 796, 3)	30
leaky_re_lu_3 (LeakyReLU)	(None, 796, 3)	0
average_pooling1d_1 (AveragePooling1D)	(None, 398, 3)	0
conv1d_4 (Conv1D)	(None, 396, 3)	30
leaky_re_lu_4 (LeakyReLU)	(None, 396, 3)	0
global_average_pooling1d (GlobalAveragePooling1D)	(None, 3)	0
dense (Dense)	(None, 1)	4
=====		
Total params: 346 (1.35 KB)		
Trainable params: 340 (1.33 KB)		
Non-trainable params: 6 (24.00 Byte)		

In [67]:

```
1 # Initialize GroupKFold with the default number of splits
2 gkf = GroupKFold()
3
4 # Initialize lists to store training and validation accuracies for each fold
5 train_accuracies = []
6 val_accuracies = []
7
8 # Initialize variables to track the best training and validation accuracies
9 best_train_accuracy = 0.0
10 best_val_accuracy = 0.0
11
12 # Iterate through the training and validation splits defined by GroupKFold
13 for train_index, val_index in gkf.split(data_array, label_array, groups=group_labels):
14     # Split data into training and validation sets
15     train_features, train_labels = data_array[train_index], label_array[train_index]
16     val_features, val_labels = data_array[val_index], label_array[val_index]
17
18     # Transpose the features if necessary
19     train_features = train_features.transpose(0, 2, 1)
20     val_features = val_features.transpose(0, 2, 1)
21
22     # Standardize features using the same scaler for both training and validation
23     scaler = StandardScaler()
24     train_features = scaler.fit_transform(train_features.reshape(-1, train_features.shape[1]))
25     val_features = scaler.transform(val_features.reshape(-1, val_features.shape[1]))
26
27     # Create and train the CNN model
28     model = cnnmodel()
29     history = model.fit(train_features, train_labels, epochs=20, batch_size=32)
30
31     # Append training and validation accuracy values to the lists
32     train_accuracies.append(history.history['accuracy'])
33     val_accuracies.append(history.history['val_accuracy'])
34
35     # Evaluate accuracies for the current fold
36     current_train_accuracy = history.history['accuracy'][-1]
37     current_val_accuracy = history.history['val_accuracy'][-1]
38
39     # Update the best training accuracy if the current training accuracy is higher
40     if current_train_accuracy > best_train_accuracy:
41         best_train_accuracy = current_train_accuracy
42
43     # Update the best validation accuracy if the current validation accuracy is higher
44     if current_val_accuracy > best_val_accuracy:
45         best_val_accuracy = current_val_accuracy
46
47 # Print the best training and validation accuracies
48 print("Best Training Accuracy:", best_train_accuracy)
49 print("Best Validation Accuracy:", best_val_accuracy)
```

2/15/23, 5:06 PM

EEG classification_Final - Jupyter Notebook

```
14/14 [=====] - 1s 99ms/step - loss: 0.6078 - accuracy: 0.5740 - val_loss: 0.6475 - val_accuracy: 0.4148
Epoch 8/20
14/14 [=====] - 1s 104ms/step - loss: 0.5777 - accuracy: 0.6059 - val_loss: 0.6036 - val_accuracy: 0.7975
Epoch 9/20
14/14 [=====] - 1s 100ms/step - loss: 0.5490 - accuracy: 0.7675 - val_loss: 0.5590 - val_accuracy: 0.8494
Epoch 10/20
14/14 [=====] - 1s 99ms/step - loss: 0.5161 - accuracy: 0.7864 - val_loss: 0.5211 - val_accuracy: 0.8469
Epoch 11/20
14/14 [=====] - 1s 100ms/step - loss: 0.4818 - accuracy: 0.8047 - val_loss: 0.4929 - val_accuracy: 0.8296
Epoch 12/20
14/14 [=====] - 1s 101ms/step - loss: 0.4612 - accuracy: 0.8024 - val_loss: 0.4716 - val_accuracy: 0.8247
Epoch 13/20
14/14 [=====] - 1s 102ms/step - loss: 0.4412 - accuracy: 0.8101 - val_loss: 0.4678 - val_accuracy: 0.8543
```

In [68]:

```
1 # Print the best training and validation accuracies
2 print("Best Training Accuracy:", best_train_accuracy)
3 print("Best Validation Accuracy:", best_val_accuracy)
```

```
Best Training Accuracy: 0.8384615182876587
Best Validation Accuracy: 0.8604061007499695
```

Model Evaluation and Performance:

The evaluation phase assesses the model's performance on the validation set using metrics such as accuracy, precision, recall, and F1-score. The goal of hyperparameter tuning is to improve the model's performance even more. Finally, the model is tested on the designated test set to ensure that it can generalize to new data.

```

1 best_clf = gscv.best_estimator_
2
3 y_pred = cross_val_predict(best_clf, features, label_array, groups=group_array)
4
5 cm = confusion_matrix(label_array, y_pred)
6
7 precision = precision_score(label_array, y_pred)
8 recall = recall_score(label_array, y_pred)
9 f1 = f1_score(label_array, y_pred)
10
11 print("Confusion Matrix:")
12 print(cm)
13 print("\nPrecision:", precision)
14 print("Recall:", recall)
15 print("F1 Score:", f1)
16 plt.figure(figsize=(8, 6))
17 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No Seizure', 'Seizure'], yticklabels=['No Seizure', 'Seizure'])
18 plt.title('Confusion Matrix')
19 plt.xlabel('Predicted Label')
20 plt.ylabel('True Label')
21 plt.show()

```

Confusion Matrix:

True Positive (TP): 760

True Negative (TN): 846

False Positive (FP): 292

False Negative (FN): 197

Precision:

Precision measures the accuracy of positive predictions.

Precision = 0.722 (72.2%)

Indicates the proportion of correctly identified seizures among all predicted seizures.

Recall (Sensitivity):

Recall measures the model's ability to capture actual positive instances.

Recall = 0.794 (79.4%)


Indicates the proportion of correctly identified seizures among all actual seizures.

F1 Score:

F1 Score is the harmonic mean of precision and recall, providing a balanced performance metric.

F1 Score = 0.757 (75.7%)

Confusion Matrix Heatmap:



The heatmap visually represents the confusion matrix.


Clearly shows the distribution of true positive, true negative, false positive, and false negative predictions.

These results collectively offer a comprehensive assessment of the model's ability to correctly classify seizures and non-seizures. The high recall indicates effective identification of seizures, while precision ensures accurate positive predictions. The F1 score provides a balanced measure, considering both precision and recall. The confusion matrix and heatmap enhance the understanding of the model's performance across different prediction outcomes.

Results and Visualization:

The project concludes with results and visualization, in which EEG data and model predictions are represented visually through plots and graphs. This comprehensive approach to EEG classification not only advances diagnostic tools in neuroscience


and medicine, but it also provides a comprehensive demonstration of the data analysis pipeline in action.



The depicted graph presents the results of a grid search over the regularization parameter C for a linear support vector machine (SVM) model. The x-axis represents different values of C , which influences the regularization strength in the SVM model. The y-axis shows the mean test scores, reflecting the performance of the SVM model under different regularization strengths.

From the graph, several key findings can be observed. Firstly, there is an optimal range of C values where the mean test score is maximized, indicating the model's peak performance. Secondly, the peak in the graph can suggest the sensitivity of the model to changes in the regularization parameter. A steep rise or fall in mean test scores indicates the model's responsiveness to adjustments in C .

Training and Validation accuracies



The training and validation accuracies for each fold, as well as the mean accuracy, provide a comprehensive overview of the model's performance across different epochs. Here are the key observations:

Training Accuracies:

The training accuracies show how well the model fits the training data over epochs.

For Fold 1, there is a consistent accuracy of approximately 59.77% throughout epochs, suggesting stable model convergence.

Folds 2, 3, and 4 exhibit varying degrees of accuracy improvement, reaching up to 78.81% for Fold 3.

Validation Accuracies:

Validation accuracies reflect the model's ability to generalize to unseen data.

Fold 1 has a constant validation accuracy of 33.33%, indicating potential underfitting or insufficient model capacity.

Folds 2 and 3 show increasing validation accuracy, peaking at 72.69% and 79.92%, respectively.

Fold 4 starts with a validation accuracy of 63.45%, demonstrating effective generalization.

Mean Accuracy:

The mean accuracy across folds is approximately 54.48%, providing an overall assessment of model performance.

This metric considers the model's ability to perform well on average across diverse datasets.

The model's performance varies across folds, suggesting potential sensitivity to data partitions. While some folds exhibit robust learning and generalization, others may require further tuning to enhance performance. Monitoring training and validation accuracies over epochs aids in identifying overfitting or underfitting issues, guiding adjustments to improve the model's overall effectiveness.

Fine-tuning hyperparameters or exploring different model architectures could be beneficial in achieving more consistent and higher accuracies across all folds.

Summary and Conclusions:

This project aimed to develop a robust EEG classification model for diagnosing neurological disorders, particularly epilepsy, leveraging advanced machine learning techniques. Through a comprehensive analysis pipeline, involving data preprocessing, feature extraction, model selection, and evaluation, the study advanced our understanding of EEG data interpretation and its potential applications in medical diagnostics.

Findings and Contributions:

- Model Performance Evaluation:**

The evaluation phase provided insightful metrics, including accuracy, precision, recall, and F1 score. The confusion matrix and heatmap offered a nuanced view of the model's ability to distinguish between seizures and non-seizures.

Precision, recall, and F1 score collectively indicated the model's effectiveness in positive prediction, capturing actual positive instances, and providing a balanced performance metric.

- Grid Search Results for SVM:**

The graph resulting from the grid search over the regularization parameter C for a linear support vector machine (SVM) model revealed important insights.

An optimal range of C values was identified where the mean test score was maximized, highlighting the model's peak performance.

The graph demonstrated the sensitivity of the model to changes in the regularization parameter, aiding in understanding the model's responsiveness to adjustments.

- **Training and Validation Accuracies Over Epochs:**

The analysis of training and validation accuracies over epochs showcased the model's learning behavior.

While some folds exhibited stable convergence, others demonstrated varying degrees of improvement, indicating potential sensitivity to data partitions.

Monitoring accuracy trends provided crucial insights into overfitting or underfitting, guiding adjustments for enhanced model effectiveness.

Recommendations and Future Directions:

- **Hyperparameter Tuning:**

Fine-tuning hyperparameters, especially for folds showing less stability, could enhance overall model performance.

Exploring different regularization strengths and model architectures may contribute to more consistent and higher accuracies.

- **Dataset Considerations:**

Further analysis of dataset characteristics and potential biases could help in understanding variations in model performance across folds.

The integration of additional diverse datasets may contribute to a more comprehensive understanding of EEG classification.

- **Interpretability and Visualization:**

Enhancements in model interpretability, such as feature importance visualization, could provide deeper insights into the model's decision-making process.

Visualizations of EEG data and predictions contribute to clearer communication of findings, especially in a medical context.

This project represents a significant step forward in the development of EEG classification models for neurological diagnosis. The findings contribute not only to the advancement of diagnostic tools in neuroscience and medicine but also serve as a comprehensive demonstration of the data analysis pipeline. Future work should focus on refining the model, exploring diverse datasets, and ensuring the interpretability of results for practical clinical applications.